Title here

Names here

Date

1 Description

2 User Manual

Our approximation of aixi is written in C++ and requires g++ for compilation.

2.1 Setup

Extract:
tar -zxvf aixi.tgz

Compile:
cd aixi
make

Run:
./aixi environment.conf

2.2 Configuration Options

Each configuration file specifies a particular environment and a set of options.
Available environments are:

- biased_rock_paper_scissor
- coinflip
- kuhn_poker
- pacman
- tiger

Other options specify parameters for the environment and the aixi agent's learning. TODO: so far I've only included options not explained in the assignment specification.

Domain	CTW depth	m	ϵ	γ	ρUCT Simulations
Biased Rock-Paper-Scissor					
Coinflip					
Kuhn-Poker					
Pacman					
Tiger					

Figure 1: Agent configurations

- mc-timelimit: The number of MC simulations per cycle.
- load-ct: Specifies a (trained) CTW for the agent to load at initialisation.
- write-ct: Write CTW to file before agent termination.

3 Experimental Results

3.1 Experimental Setup

TODO:

- List configurations used for each environment.
- List hardware (cpu/clock speed/cache/ram)

3.2 Discussion

How well did it do.

Include results to do with forgetting past model - changing environments

Include statistics about cycles required for optimal performance, time per cycle as in the VNHS paper [1].

TODO: Anyone have a bibtex library already? If no one does, it may be easier to just do references manually.

References

[1] J. Veness et al. Reinforcement learning via AIXI approximation Technical report, Australian National University, 2009.